

1.2 Regularized linear regression cost function

$$J(\theta) = \frac{1}{2m} \left(\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 \right) + \frac{\lambda}{2m} \left(\sum_{j=1}^n \theta_j^2 \right),$$

```
1 J = sum((X * theta - y).^2)/2/m;  
2 temp = theta;  
3 temp(1) = 0;  
4 J += lambda * sum(temp.^2)/2/m;
```

1.3 Regularized linear regression gradient

$$\frac{\partial J(\theta)}{\partial \theta_0} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \quad \text{for } j = 0$$
$$\frac{\partial J(\theta)}{\partial \theta_j} = \left(\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \right) + \frac{\lambda}{m} \theta_j \quad \text{for } j \geq 1$$

```
1 grad(0) = sum((X*theta-y).*X(:,1))/m;  
2 for i=2:m,  
3   grad(i) = sum((X * theta - y) .* X(:,i))/m + lambda*theta(i)/m;  
4 end
```

2 Bias-variance

2.1 Learning curves

You can use the `trainLinearReg` function to find the θ parameters. Note that the `lambda` is passed as a parameter to the `learningCurve` function. After learning the θ parameters, you should compute the **error** on the training and cross validation sets. Recall that the training error for a dataset is defined as

$$J_{\text{train}}(\theta) = \frac{1}{2m} \left[\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 \right].$$

```
1 for i=1:m,  
2   theta = trainLinearReg(X(1:i,:),y(1:i),lambda);  
3   error_train(i) = linearRegCostFunction(X(1:i,:),y(1:i),theta,0);  
4   error_test(i) = linearRegCostFunction(Xval,yval,theta,0);  
5 end
```

3 Polynomial regression

polyFeatures:

Now, you will add more features using the higher powers of the existing feature x in the dataset. Your task in this part is to complete the code in `polyFeatures.m` so that the function maps the original training set X of size $m \times 1$ into its higher powers. Specifically, when a training set X of size $m \times 1$ is passed into the function, the function should return a $m \times p$ matrix `X_poly`,

where column 1 holds the original values of X , column 2 holds the values of $X.^2$, column 3 holds the values of $X.^3$, and so on. Note that you don't have to account for the zero-th power in this function.

```
1 for i=1:p,  
2   X_poly(:,i) = X.^i;  
3 end
```

3.3 Selecting lambda using a cross validation set

```
1 lambda_vec = [0 0.001 0.003 0.01 0.03 0.1 0.3 1 3 10];  
2 for i = 1:length(lambda_vec),  
3   lambda = lambda_vec(i);  
4   theta = trainLinearReg(X,y,lambda);  
5   error_train(i) = linearRegCostFunction(X,y,theta,0);  
6   error_test(i) = linearRegCostFunction(Xval,yval,theta,0);  
7 end
```